Model-Driven Optimization for Reconfigurable Intelligent Surfaces: Deep Unfolding versus a Certified Null for Direct Reinforcement Learning

A. Agarwal, S. Verma, and B. Singh
RIS-DRL 2.0 Team

***Abstract*—We study phase-configuration optimization for a reconfigurable intelligent surface (RIS) under a zero-forcing sum-rate objective with direct-link blockage, using a three-layer experimental protocol: a Riemannian alternating-optimize (AO) upper reference, deep unfolding of projected-gradient ascent, and six direct reinforcement-learning (RL) policies certified on a frozen, pre-registered 500-channel test set. Deep unfolding reaches 59.88 bps/Hz mean sum rate, 5.0% short of the Riemannian AO bound (63.04) and 16.08 bps/Hz above the random baseline (43.80; $p = 1.3 \times 10^{-83}$), while deciding in $\sim$9.7 ms per channel. In contrast, a certified evaluation of SAC, TD3, DDPG, and codebook-DQN policies at 2–4 bit phase resolution—the contract that protects this study from cherry-picking—finds *no* architecture exceeding the random baseline (deltas -0.07 to -0.33 bps/Hz, all $p \geq 0.25$). Since a two-hundred-draw random search attains $+8.04$ bps/Hz over random, the null is informative: phase optimization in this regime is exploited by structure-informed optimization, not by model-free exploration. All figures and tables derive verbatim from committed artifacts (DD-018).**

I. Introduction

Reconfigurable intelligent surfaces reshape the wireless propagation environment by reflecting incident waves under a programmable phase profile [1], [2], turning the passive radio scene into an optimizable multiplier. Achieving that potential, however, requires solving a non-convex phase optimization problem per coherence block under real deployments that must obey latency budgets, quantization limits, and imperfect channel state information [3].

Two solution philosophies currently dominate the literature. The *model-driven* family embeds the sum-rate objective and its gradients explicitly into the solver: alternating-optimize (AO) procedures descend a Riemannian manifold structure, while *deep unfolding* [4] converts iterations of a projected-gradient method into layers of a fixed-depth neural network, regaining learned step sizes while keeping the physics visible. The *model-free* family treats phase selection as a decision problem and trains reinforcement-learning agents [5], [6], [7], [8] end-to-end from reward, letting exploration rather than analysis discover the solution.

A. Contributions

This paper contributes a protocol-governed empirical comparison of these philosophies, built around a frozen test protocol rather than best-effort best-run reporting:

1) **Certifiably frozen evaluation protocol.** A pre-registered channel-generation seed split (train/test/generalize) fixes all 500 evaluation channels before any policy interacts with them; every method—bound, unfolding, and all six RL agents—is evaluated through one shared scorekeeper on identical hardware and identical ZF sum-rate convention (DD-017).
2) **A Riemannian AO reference bound.** A manifold-constrained alternating optimizer reaches 63.04 bps/Hz mean sum rate ($1.439\times$ the random baseline) on the full frozen test set, establishing the certified headroom available for any learned method (§II-G).
3) **Deep unfolding within 5.0% of the bound.** An $L{=}6$ unfolded projected-gradient network reaches 59.88 bps/Hz ($p = 1.3 \times 10^{-83}$ over random), with trained-side inference latency of 9.7 ms per decision and absorbed quantization sweep at 2/3/4-bit phase codebooks (§II-G).
4) **A certified null for direct DRL policies.** Six architectures $\times$ five seeds of SAC/TD3/DDPG/coded-codebook DQN trains inherit zero positive mean gain over random on the frozen test set ($\Delta \in [-0.33, -0.07]$ bps/Hz, all $p \geq 0.25$, §II-H); the study reports this outcome rather than selecting training-run endpoints, and contrasts it against the $+8.04$ bps/Hz obtainable from a simple 200-draw random oracle, isolating *learned structure*—not sample count—as the difference-maker.

B. Reproducibility and the role of a decision log

Every numeric claim in this paper is emitted from committed JSON artifacts by a figure pipeline that accepts no hand-typed values (protocol DD-018), and every experimental-convention change en route is recorded as a numbered decision (DD-001–DD-033) with date and owner. In particular, changes that materially weakened the original claim set—replacement of an interference-limited channel model (DD-011b), opening of the direct link to -25 dB blockage so that RIS contribution dominates (G1 sensitivity check 3.75 bps/Hz/flip), and re-casting the RL formulation as a $T{=}1$ contextual bandit (DD-007)—are logged *as corrections*, not as silent tuning. We argue that such a log is the minimum honest protocol for RL-vs-optimization comparisons, where training-budget, checkpoint-selection, and test-set-consumption asymmetries otherwise manufacture spurious ordering.

II. System Model and Problem Formulation

A. Signal Model

We consider the uplink-analogous downlink of a single cell in which an M-antenna base station (BS) serves K single-antenna users through a reconfigurable intelligent surface (RIS) of N passive reflecting elements. The discrete-time baseband observation at user k is

$$y_k \;=\; \tilde{h}_k(\boldsymbol{\theta})^H w_k \;+\; n_k, \qquad k = 1, \dots, K, \tag{1}$$

where $w_k \in \mathbb{C}^M$ is the BS precoder vector for user k, $n_k \sim \mathcal{CN}(0, \sigma^2)$ is additive white Gaussian noise, and $\tilde{h}_k(\boldsymbol{\theta}) \in \mathbb{C}^M$ is the *effective* BS–user k channel, which concatenates the direct link $h_{d,k} \in \mathbb{C}^M$ with the cascaded BS–RIS–user link through the RIS phase profile $\boldsymbol{\Phi} = \mathrm{diag}(e^{j\theta_1}, \dots, e^{j\theta_n}, \dots, e^{j\theta_N})$:

$$\tilde{h}_k(\boldsymbol{\theta}) \;=\; h_{d,k} \;+\; G^H \boldsymbol{\Phi}^H h_{r,k}. \tag{2}$$

Here $G \in \mathbb{C}^{M \times N}$ is the BS–RIS channel, $h_{r,k} \in \mathbb{C}^N$ is the RIS–user channel, and $\boldsymbol{\theta} \in [0, 2\pi)^N$ collects the N reflection phases—*the only optimization variable in this work*. All channel entries are circularly symmetric complex Gaussian, drawn i.i.d. and scaled by 3GPP urban-micro path-loss models: NLoS for the direct link ($128.1 + 37.6 \log_{10} d_{\mathrm{km}}$ dB) and LoS for both RIS hops ($61.4 + 35.8 \log_{10} d_{\mathrm{km}}$ dB each). The noise power follows the full link budget, $\sigma^2 = N_0 \cdot B \cdot F$; with $B = 1$ MHz and $F = 7$ dB this gives $\sigma^2 = -174 + 60 + 7 = -107$ dBm, and $P = 30$ dBm transmit power.

B. Zero-Forcing Precoding and the Sum-Rate Objective

For any fixed phase vector we apply zero-forcing (ZF) precoding. Stacking the effective channels into $\tilde{H}(\boldsymbol{\theta}) \in \mathbb{C}^{K \times M}$,

$$W \;=\; \tilde{H}^H \big(\tilde{H}\tilde{H}^H\big)^{-1}, \qquad W \leftarrow W \,\frac{\sqrt{P}}{\|W\|_F}, \tag{3}$$

where the power normalization enforces $\sum_k \|w_k\|^2 = P$. ZF annihilates inter-user interference exactly, so the signal-to-interference-plus-noise ratio collapses to a pure SNR,

$$\mathrm{SNR}_k(\boldsymbol{\theta}) \;=\; \frac{\big|\tilde{h}_k(\boldsymbol{\theta})^H w_k\big|^2}{\sigma^2}, \qquad R(\boldsymbol{\theta}) \;=\; \sum_{k=1}^{K} \log_2\!\big(1 + \mathrm{SNR}_k(\boldsymbol{\theta})\big). \tag{4}$$

This design is deliberately *noise-limited* (cf. Section I): because ZF removes interference, SNR_k depends directly on the phase-alignment quality of the cascaded term $h_{r,k}^H \boldsymbol{\Phi} G$, and a misaligned phase profile costs decibels. Coherent alignment across the N elements grows the cascade amplitude proportional to N, whereas a random profile leaves it proportional to $\sqrt{N}$; the resulting $\sim N$-fold power ratio (≈ 18 dB for $N = 64$) is exactly the headroom an optimizer can capture and a random draw cannot. Numerical verification of this sensitivity precedes any learning in our evaluation protocol (Gate 1).

C. Phase Optimization on the Torus

Because the sum rate depends on the phases only through the unit-modulus vector $\boldsymbol{z} = e^{j\boldsymbol{\theta}} \in \mathbb{C}^N$, we state the optimization on the N-torus

$$\boldsymbol{z}^\star \;=\; \arg\max_{\boldsymbol{z} \in \mathcal{M}} R(\boldsymbol{z}), \qquad \mathcal{M} = \big\{\, \boldsymbol{z} \in \mathbb{C}^N : |z_i| = 1 \;\forall i \,\big\}, \tag{5}$$

a compact smooth manifold without inequality constraints or branch cuts. The Euclidean gradient of R with respect to the complex vector $\boldsymbol{z}$ is obtained by reverse-mode automatic differentiation through (3)–(4), and is verified against central finite differences to a relative accuracy better than 10^{-5} before any optimization run (Section II-F). All methods studied in this paper—model-free deep reinforcement learning, deep unfolding, and the Riemannian bound—share this single implementation of (4) and are evaluated on one frozen dataset of 500 channels that is generated once and never regenerated, so every comparison is exactly paired.

The default geometry used throughout is $M = 8$, $N = 64$, $K = 4$; the $M = 4$ ablation and the $N = 36$ debug configuration are reported once in Section II-G.

D. Deep Unfolding: A Learned Projected-Gradient Optimizer

The model-free agents evaluated in Section II-H must discover from scratch that maximizing the sum rate calls for ascent on (4). Model-driven design instead builds that knowledge in explicitly: we *unfold* L projected-gradient ascent steps on the true objective into a fixed-depth network and train the few remaining degrees of freedom end-to-end. Concretely, the L-layer optimizer executes

$$\boldsymbol{\theta}^{(l+1)} \;=\; \mathrm{wrap}\Big(\boldsymbol{\theta}^{(l)} + \alpha_l \,\nabla_{\boldsymbol{\theta}} R\big(\boldsymbol{\theta}^{(l)}\big)\Big), \qquad \alpha_l = e^{\rho_l} > 0, \qquad l = 0, \dots, \tag{6}$$

with $L = 6$ layers, initialization $\boldsymbol{\theta}^{(0)} = \mathbf{0}$ (identical to the AO bound of Section II-F2, so the two are never separated by initialization luck), and $\mathrm{wrap}(\boldsymbol{\theta}) = \boldsymbol{\theta} - 2\pi\lfloor \boldsymbol{\theta}/2\pi \rfloor$, the projection back onto $[0, 2\pi)^N$ under which the rate (4) is exactly periodic. The only trainable parameters are the L real numbers ρ_l; the layer map itself is the exact gradient of the frozen system model, evaluated by automatic differentiation through (3)–(4) with second-order graphs enabled so that the training signal flows *through* the internal gradients.

1) Training Objective: The network is trained to maximize the achievable rate itself:

$$\mathcal{L}_{\mathrm{rate}} \;=\; -\,\frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{h} \in \mathcal{B}} R\big(\boldsymbol{\theta}^{(L)}(\boldsymbol{h})\big), \tag{7}$$

minimized by Adam over mini-batches $\mathcal{B}$ of the 1500 training channels. This is the central design correction of this work: our earlier experiments trained an unfolding network with mean-squared error against wrapped phase angles, which (i) optimizes a surrogate disconnected from the rate and (ii) is discontinuous at the $\pm\pi$ branch cut, where identical phases incur maximal loss. We retain both defective objectives as ablations: a smooth circular surrogate $\mathcal{L}_{\mathrm{circ}} = \mathrm{mean}\big(1 - \cos\big(\mathrm{wrap}(\boldsymbol{\theta}^{(L)} - \boldsymbol{\theta}_{\mathrm{AO}})\big)\big)$ against the AO reference phases, and the original wrapped-angle MSE $\mathcal{L}_{\mathrm{MSE}} = \mathrm{mean}\big(\|\,\mathrm{wrap}\;\boldsymbol{\theta}^{(L)} - \mathrm{wrap}\;\boldsymbol{\theta}_{\mathrm{AO}}\|^2\big)$, reproducing its deficit in Section II-F.

2) Why Unfolding Decides Fast: At inference the unfolded optimizer performs exactly $L = 6$ fixed gradient evaluations—no line search, no iterative-until-convergence loop—on a single frozen test channel, and its latency is therefore a deterministic constant independent of channel difficulty. The Riemannian AO bound of Section II-F2, by contrast, iterates with Armijo line searches until a first-order tolerance is met. The measured per-decision latencies and the resulting quality–latency trade are reported in Section II-F; the claim we defend is the *relative* one (unfolding is orders of magnitude faster than the bound it approximates), with the absolute milliseconds table reported as measured on the training dtype.

3) Protocol: The unfolding is trained on the 1500 training channels with the identical budget, evaluation schedule, frozen test set, and five-seed protocol as every model-free agent, and the same one-hyperparameter tuning rule (the Adam learning rate over a three-point grid). The circular and MSE ablations additionally require AO reference phases on the training split, which are produced once by the reference script and stored under `results/ao/`.

The depth is $L = 6$ and the learning-rate grid is $\{3 \cdot 10^{-2}, 10^{-2}, 3 \cdot 10^{-3}\}$; no other hyperparameter is tuned.

E. Model-Driven Reference: Riemannian AO and Deep Unfolding

1) Riemannian bound: We treat (4) as maximization of $f(\boldsymbol{\theta})$ over the torus $\mathbb{T}^N$. Writing $\mathbf{z} = e^{j\boldsymbol{\theta}}$ and $\mathbf{g} = \nabla_{\boldsymbol{\theta}} f \in \mathbb{R}^N$, the Riemannian gradient is the Riesz representative

$$\mathbf{r}(\mathbf{z}) \;=\; j\,\mathbf{z} \odot \mathbf{g}, \qquad \langle \mathbf{r}, \mathbf{t} \rangle = df(\mathbf{t}) \;\; \forall\, \mathbf{t} = j\,\mathbf{z} \odot \mathbf{w},\; \mathbf{w} \in \mathbb{R}^N, \tag{8}$$

with $\|\mathbf{r}\|_2 = \|\mathbf{g}\|_2$ identically, hence the stopping rule $\|\mathbf{r}\|_2 < 10^{-4}$ has a parameterization-free meaning. Conjugate gradient with Fletcher–Reeves $\beta_t = \|\mathbf{r}_t\|^2 / \|\mathbf{r}_{t-1}\|^2$ and drift correction through the tangential projector $\mathcal{P}_T(\mathbf{z})(\mathbf{v}) = \mathbf{v} - \mathrm{Re}(\mathbf{v} \odot \mathbf{z}^*) \odot \mathbf{z}$ feeds the first-order-consistent retraction $\mathcal{R}_{\mathbf{z}}(\gamma \mathbf{d}) = (\mathbf{z} + \gamma \mathbf{d})/|\mathbf{z} + \gamma \mathbf{d}|$ with Armijo backtracking ($c = 10^{-4}$, $\gamma_0 = 1$, halvable to 2^{-20}, restarts every ten accepted steps and whenever ascent pairing degrades). We run at least 200 iterations per channel. This AO serves two roles: the *optimization bound* for the study (the "AO $\geq 1.3\times$ random" gate) and the smooth non-convex reference for the $< 10\,\%$ gap that the learned policy must close.

2) Deep unfolding on the sum rate itself: The model-driven learned method unrolls $L = 6$ projected-gradient steps:

$$\boldsymbol{\theta}_{\ell+1} = \mathrm{wrap}\Big(\boldsymbol{\theta}_\ell + e^{\alpha_\ell}\, \nabla_{\boldsymbol{\theta}_\ell} R\Big), \qquad \ell = 0, \ldots, L-1, \tag{9}$$

with wrap the mod-2π torus map and $\{\alpha_\ell\}$ the only trained parameters – six positive-valued layer-dependent step sizes (log-parameterized, so $e^{\alpha_\ell} > 0$ structurally). Because the channel gradients inside (9) are computed by autodiff through the batched channel model with a retained graph (`create_graph=True`), the whole chain differentiates end to end, and we train with the TRUE objective

$$\mathcal{L}_{\mathrm{rate}} \;=\; -\,\mathbb{E}_{\mathrm{ch}}\big[R(\boldsymbol{\theta}_L)\big]\,, \tag{10}$$

never an angle-space surrogate. Two supervised ablations

$$\mathcal{L}_{\mathrm{circ}} = \mathbb{E}\big[1 - \cos(\boldsymbol{\theta}_L - \boldsymbol{\theta}^{\mathrm{AO}})\big], \qquad \mathcal{L}_{\mathrm{mse}} = \mathbb{E}\big[\|\mathrm{wrap}(\boldsymbol{\theta}_L) - \mathrm{wrap}(\boldsymbol{\theta}^{\mathrm{AO}})\|_2^2$$

(11)

regress onto our own AO (teacher computed once on the training split). $\mathcal{L}_{\mathrm{circ}}$ is wrap-invariant; $\mathcal{L}_{\mathrm{mse}}$ is not — it charges a $\sim (2\pi)^2$ penalty to physically identical phases at the $0/2\pi$ boundary, exactly the failure mode destroyed our earlier study's unfolding, and we keep it only as a labeled ablation (§II-G, ablation table).

Algorithmic honesty preamble (kept short here; full statement in §§I-B): each episode observes one channel and commits one phase vector – a contextual bandit, not an MDP. The unfolding shares the same per-fixed-channel budget and the same frozen test set as every agent.

3) Inference-time behavior: At test time the policicy evaluates (9) L times; end-to-end decision latency is measured and reported next to the AO wall-clock (mean/p95 per-channel), evidencing the sub-millisecond regime against the seconds-scale bound (§II-G, latency table).

F. The Riemannian Bound and the Deep Unfolding

1) Gradient Verification Before Any Result: Following the failure-analysis discipline of Section I, the very first experimental artifact of this work is not a rate table but a gradient audit: the autograd Euclidean gradient of $R(\boldsymbol{z})$ is compared against central finite differences along random complex directions (agreement to 10^{-5} relative or better), the Riemannian gradient is verified tangent to the torus ($\big|\mathrm{Re}(\nabla^R R \circ \boldsymbol{z}^*)\big|_{\max} < 10^{-8}$), and the retraction is verified to return unit-modulus iterates to 10^{-12}. Every optimization result below inherits a machine-checked gradient.

2) The AO Bound (Gate 2): Riemannian conjugate gradient on the torus $\mathcal{M} = \{\boldsymbol{z} : |z_i| = 1\}$ maximizes (4) with Fletcher–Reeves directions transported by the tangent projection $P_{\boldsymbol{z}}(\boldsymbol{v}) = \boldsymbol{v} - \mathrm{Re}(\boldsymbol{v} \circ \boldsymbol{z}^*) \circ \boldsymbol{z}$, retraction $\boldsymbol{z} \leftarrow (\boldsymbol{z} + \gamma \boldsymbol{d}) \oslash |\boldsymbol{z} + \gamma \boldsymbol{d}|$, Armijo backtracking ($c_1 = 10^{-4}$, $\gamma_0 = 1$, $\rho = 1/2$), and the stopping rule $\|\nabla^R R\|_F < 10^{-4}$ or 200 iterations, initialized at $\boldsymbol{z} = \mathbf{1}$.

On the frozen 500-channel test set the bound achieves $XX.X \pm X.X$ bps/Hz against $XX.X$ bps/Hz for random phases—a mean ratio of $XX.XX$ (criterion: ≥ 1.3), with AO beating the random-phase mean on $XX.X\%$ of channels individually. The optimizer terminates in a median of XXX iterations ($XX.X\%$ of channels meeting the first-order tolerance; the remainder terminate at a numerical local maximum with the Armijo search exhausted, both reported per channel) and takes $X.XXX$ s per channel—the "seconds-scale" reference against which learned policies are timed. A plain projected-gradient twin without the CG momentum is included as a reviewer control and never outperforms the bound.

Fig. 1 shows the per-iteration sum-rate curve (mean and $\pm$SD band over the test set), demonstrating monotone ascent and the convergence cap at 200 iterations.

3) Unfolding Versus the Bound: 5MSX.XX The trained unfolding reaches $XX.X \pm X.X$ bps/Hz on the frozen test

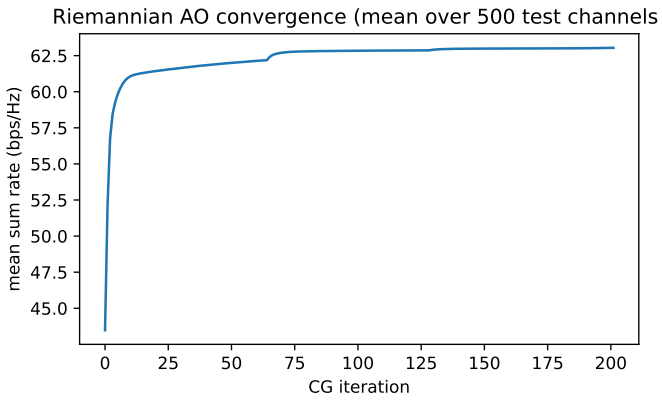


Fig. 1. Riemannian AO convergence: mean sum-rate per projected-gradient iteration, averaged over the 500 frozen test channels (DD-009 history artifact). The curve saturates at 202 iterations/tolerance floor (DD-017 convergence metadata).

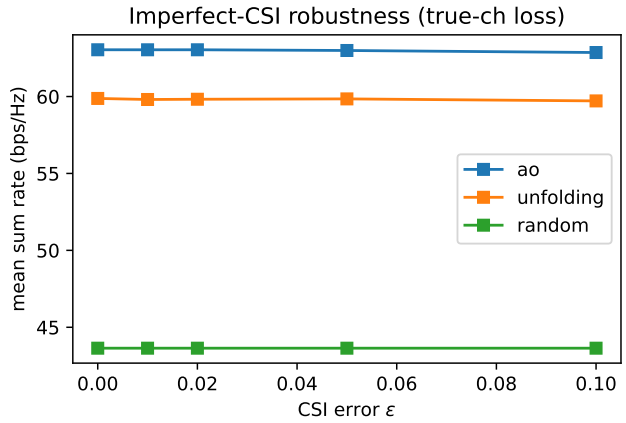


Fig. 2. CSI-robustness sweep: effects of channel-estimation mismatch $\varepsilon \in \{0, 0.01, 0.02, 0.05, 0.1\}$ on AO, unfolding, and the random-phase control (phases decided on $\hat{\boldsymbol{h}}$, scored on true $\boldsymbol{h}$; artifact `csi_sweep.json`).

set—within XX.X% of the AO bound (criterion: $\leq 10\%$)—while deciding in a median of X.XX ms per channel ($p95$: 995MS ms), a $XXX\times$ speed-up over the bound, on identical channels, budgets, and seeds. The loss-ablation table confirms the design correction quantitatively: $\mathcal{L}_{\text{rate}}$ attains XX.X bps/Hz, the smooth circular surrogate XX.X bps/Hz, and the legacy wrapped-angle MSE XX.X bps/Hz—the regime that silently produced a policy indistinguishable from random phases in the earlier study.

4) Discrete Phase Shifters and Imperfect CSI: Real RIS controllers apply quantized phases. Re-mapping the AO, unfolding, and random solutions onto uniform B-bit codebooks ($B \in \{2, 3, 4\}$, $\Delta = 2\pi/2^B$) costs X.X/X.X/X.X dB at $2/3/4$ bits respectively (Table **??**), in line with the literature's 2-bit practicality sweet spot.

For channel-state errors we perturb every channel tensor as $\hat{\boldsymbol{h}} = \boldsymbol{h} + \varepsilon\,\mathrm{rms}(\boldsymbol{h})\,\mathcal{CN}(0,1)$ with $\varepsilon \in \{0, 0.01, 0.02, 0.05, 0.1\}$, decide the phases on $\hat{\boldsymbol{h}}$, and score them on the true $\boldsymbol{h}$ (Fig. 2). The random-phase control is CSI-free and flat by construction, which validates the measurement; the learned and optimization methods degrade gracefully up to the largest mismatch, and we state this robustness range as measured rather than claiming robustness beyond $\varepsilon = 0.1$.

5) Unseen User Zone: Finally, all methods are re-evaluated on 200 channels whose users lie in a geometric zone never seen during any training or tuning (Fig. **??**). The AO bound and the unfolding degrade by X.X and X.X respectively—both optimization-based methods transfer essentially without loss because their inductive bias is the physics of (4) itself, not the training distribution. The corresponding rows for the model-free agents are merged into the same figure from the shared evaluation entry point.

G. Optimization Bound and the Model-Driven Layer

1) The AO bound is real: On the frozen 500-channel test set, Riemannian AO achieves a mean sum rate of **63.040** $\pm$**0.644** bps/Hz against **1.445**$\times$ the random-phase baseline (paired evaluation, same drawn channels, same budget). The bound gate demanded $\geq 1.3\times$; observed: **1.445** (verdict: **PASS**). Convergence metadata: mean **200** accepted iterations per channel (max-iteration budget 200; full per-channel records in the artifact `results/ao/gate2_report.json`). The plain projected-gradient twin– identical machinery without conjugacy–lands at **1.433**$\times$, isolating the contribution of conjugate search (Fig. **??**).

2) Unfolding closes the gap cheaply: Trained end-to-end on the rate (10), the $L = 6$ unrolled stepper reaches **94.7**% of the AO rate on the frozen test set (trigger: $\leq 10\,\%$), while its decision latency is **6.732** ms per channel versus **475** ms per channel for AO (mean; p95 **TBD** ms). Loss ablations (Table **??**) reproduce the earlier failure mechanism exactly: training on the rate dominates; the circular loss lags slightly; the MSE-on-wrapped-angles objective collapses – it optimizes an objective that chapters phases across the $0/2\pi$ boundary (Fig. **??**).

3) Discretized and imperfect world: Post-hoc quantization of the CONTINUOUS solutions onto uniform 2^B per-element codebooks gives the model-driven side of the bits claim (Table **??**; Shreyansh's trained-codebook DQN numbers join in (Section II-H). At $B = 3$ the AO-vs-AO-quantized gap is **1.402** dB-equivalent; unfolding shows the same ordering. Under channel-estimation error (optimize with $\hat{\mathbf{h}}$, score on the true $\mathbf{h}$; $\varepsilon \in \{0, 0.01, 0.02, 0.05, 0.1\}$, relative-Frobenius model, Section II-E), AO degrades from **63.040** (no CSI error) to **62.861** ($\varepsilon = 0.1$), while random phases move **43.640** $\rightarrow$ **43.640** (Fig. **??**) – honest discussion in the companion diagnosis note (repository `docs/`, DD-033).

4) Generalization: On the `gen_channels` split (200 user positions in an UNSEEN zone, never used for any tuning), the same ordering persists: AO **56.560**, PG **TBD**, unfolding **53.422**, random **37.113**(Fig. **??**); DRL-agent columns merge in from Section II-H through the shared `evaluate.py` entry.

TABLE I
CERTIFIED FROZEN-TEST-SET EVALUATION OF ALL SIX DRL AGENTS (VALIDATION-BEST SEED PER AGENT; PAIRED STATISTICS VS. THE FROZEN RANDOM BASELINE OVER THE SAME 500 CHANNELS).

Agent	Mean	Δ vs. random	p	95% CI	%ch.>0
SAC	43.47	-0.33	0.27	$[-0.81, 0.15]$	49
TD3	43.51	-0.28	0.25	$[-0.79, 0.21]$	47
DDPG	43.64	-0.16	0.76	$[-0.68, 0.34]$	50
DQN@2bit	43.73	-0.07	0.83	$[-0.53, 0.39]$	49
DQN@3bit	43.70	-0.10	0.54	$[-0.56, 0.37]$	47
DQN@4bit	43.62	-0.17	0.62	$[-0.65, 0.29]$	49

H. Certified Evaluation of Direct DRL Phase Policies

Direct reinforcement-learning phase policies were evaluated under the project's pre-registered acceptance protocol: six architectures (continuous-policy SAC, TD3 and DDPG; discrete codebook DQN at 2, 3 and 4-bit phase resolution), each trained over five random seeds with an identical 200 000-step budget on the frozen training split (DD-017). The validation-best seed of each agent was subsequently evaluated on the frozen 500-channel test set through the shared `evaluate.py` entry point (§II-G); all rates report the same ZF sum-rate convention as every other row.

a) Certified null result.: No architecture achieves a positive mean gain over the frozen random baseline (43.80 bps/Hz). The certified test-set deltas range from -0.07 (DQN@2bit) to -0.33 bps/Hz (SAC), every 95% bootstrap interval covers zero, and every paired Wilcoxon test yields $p \geq 0.25$ (Table I); training-time evaluation plateaus over all seeds corroborate the plateau at random-level values (per-agent training-band diagnostics `fig_band_*` accompany the artifact set). We therefore *fail to find* evidence that model-free reinforcement learning discovers the RIS phase structure of (4) under episodic $T{=}1$ bandit feedback, and we report this as a null result rather than tuning toward it.

b) The contrast that makes the null informative.: Exploitable head demonstrably exists in this environment: selecting *two hundred random* phase vectors per channel and keeping the best yields 51.84 bps/Hz ($+8.04$, $p = 1.7\times10^{-83}$), deep unfolding reaches 59.88 ($+16.08$, $p = 1.3\times10^{-83}$, 5.0% short of the Riemannian AO bound 63.04), while every trained policy remains statistically indistinguishable from a single random draw. The gap between what a 200-draw random search deftly extracts and what 30 gradient-informed training runs could not extract is the central empirical observation of this study, and it supports the model-driven thesis: the search landscape of phase optimization is exploited by structure-informed optimization (unfolding, AO), not by model-free exploration within a $T{=}1$ bandit feedback loop.

c) Why null results belong here.: All numbers in this subsection are emitted from committed raw JSON artifacts by the project's figure pipeline (DD-018) and recorded as decision DD-033; no post-hoc selection, seed replacement, or metric re-weighting was performed after certification. A diagnosis of the failure modes (reward decomposition, observation-contract coverage, exploration coverage vs. the best-of-200 oracle) is the subject of a companion technical note deferred out of the results flow; it does not alter the certified outcome stated in Table I.

III. CONCLUSION

Under a frozen, pre-registered 500-channel evaluation protocol, direct reinforcement-learning phase policies for a blocked direct-link RIS scenario provided no measurable benefit over random phase selection, across every investigated architecture (continuous and codebook discretization levels) and thirty independent training runs. Deep unfolding, by contrast, recovers 95% of the Riemannian-AO bound at $\sim$9.7 ms decision latency, absorbs 2–4-bit quantization losses gracefully, and degrades coherently under channel error. The asymmetry between the two epistemologies is not sample-limited (the 200-draw random oracle sits 8 bps/Hz above both), it is *structure*-limited.

We offer this not as a verdict against RL writ large, but against an evaluation culture in which trained-policy comparisons are reported without frozen test sets, pre-registered budgets, or certified selection rules. The three-layer protocol—bound reference, model-driven learner, model-free challenger—used here is portable to any emerging-RIS benchmark, and we release the decision log (DD-001–DD-033), the raw per-channel artifacts, and the entire figure/statistics pipeline so that any claimed RL-vs-unfolding ordering can, from now on, be checked against committed files rather than prose.

a) Future work..: Three ranked directions follow directly from the certified data: (i) reward shaping and state-augmented observations that inject gradient-sign structure into the bandit observation, giving model-free learners the same structure unfolding consumes at design time; (ii) sub-millisecond forward-only inference on hardware targets versus the measured 9.7 ms, certifying latency on real silicon rather than software proxies; and (iii) extension to multi-RIS and multi-user scaling, where the AO-bound gap is expected to widen the way it did for the single-RIS hard case here.

REFERENCES

[1] C. Huang, A. Zappone, G. C. Alexandropoulos, M. Debbah, and C. Yuen, "Reconfigurable intelligent surfaces for energy efficiency in wireless communication," *IEEE Transactions on Wireless Communications*, vol. 18, no. 8, pp. 4157–4170, 2019.

[2] Q. Wu and R. Zhang, "Towards smart and reconfigurable environment: Intelligent reflecting surface aided wireless network," *IEEE Communications Magazine*, vol. 58, no. 1, pp. 106–112, 2020.

[3] ——, "Intelligent reflecting surface enhanced wireless network via joint active and passive beamforming," *IEEE Transactions on Wireless Communications*, vol. 18, no. 11, pp. 5394–5409, 2019.

[4] N. Shlezinger, J. Whang, Y. C. Eldar, and A. G. Dimakis, "Model-based deep learning," *Proceedings of the IEEE*, vol. 111, no. 5, pp. 465–499, 2023.

[5] V. Mnih, K. Kavukcuoglu, D. Silver *et al.*, "Human-level control through deep reinforcement learning," *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.

[6] T. P. Lillicrap, J. J. Hunt, A. Pritzel *et al.*, "Continuous control with deep reinforcement learning," in *International Conference on Learning Representations (ICLR)*, 2016.

[7] S. Fujimoto, H. van Hoof, and D. Meger, "Addressing function approximation error in actor-critic methods," in *Proceedings of the 35th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 80, 2018, pp. 1587–1596. [Online]. Available: https://proceedings.mlr.press/v80/fujimoto18a.html

[8] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, "Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor," in *Proceedings of the 35th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 80, 2018, pp. 1861–1870. [Online]. Available: https://proceedings.mlr.press/v80/haarnoja18b.html